
Byte Ordering and Byte Manipulation Functions

Byte Ordering Functions

- Unfortunately, not all computers store the bytes that comprise a multibyte value in the same order.
- Consider a 16-bit integer that is made up of 2 bytes. There are two ways to store this value.
 - **Little Endian** – In this scheme, low-order byte is stored on the starting address (A) and high-order byte is stored on the next address ($A + 1$).
 - **Big Endian** – In this scheme, high-order byte is stored on the starting address (A) and low-order byte is stored on the next address ($A + 1$).

Byte Ordering Functions

- To allow machines with different byte order conventions communicate with each other, the Internet protocols specify a canonical byte order convention for data transmitted over the network. This is known as Network Byte Order.
- While establishing an Internet socket connection, you must make sure that the data in the `sin_port` and `sin_addr` members of the `sockaddr_in` structure are represented in Network Byte Order.

Byte Ordering Functions

- Functions for converting data between a host's internal representation and Network Byte Order are as follows –
- **unsigned short htons(unsigned short hostshort)** – This function converts 16-bit (2-byte) quantities from host byte order to network byte order.
- **unsigned long htonl(unsigned long hostlong)** – This function converts 32-bit (4-byte) quantities from host byte order to network byte order.
- **unsigned short ntohs(unsigned short netshort)** – This function converts 16-bit (2-byte) quantities from network byte order to host byte order.
- **unsigned long ntohl(unsigned long netlong)** – This function converts 32-bit quantities from network byte order to host byte order.

Byte Ordering Functions

- These functions are macros and result in the insertion of conversion source code into the calling program.
 - On little-endian machines, the code will change the values around to network byte order.
 - On big-endian machines, no code is inserted since none is needed; the functions are defined as null.

Host : Big/Little Endian
Network: Big Endian

Program to Determine Host Byte Order

```
#include <stdio.h>

int main(int argc, char **argv)
{
    union {
        short s;
        char c[sizeof(short)];
    }un;

    un.s = 0x0102; //hexadecimal
    if (sizeof(short) == 2)
    {
        if (un.c[0] == 1 && un.c[1] == 2)
            printf("big-endian\n");
        else if (un.c[0] == 2 && un.c[1] == 1)
            printf("little-endian\n");
        else printf("unknown\n");
    }
}
```

```
else {
    printf("sizeof(short) = %d\n",
        sizeof(short));
}
exit(0);
}
```

- The elements of a union occupy a common "piece" of memory. So un.s and un.c[] refer to same memory
- s= 0x0102, therefore
- If c={0x02, 0x01} → Little Endian
- If c={0x01, 0x02} → Big Endian

Byte Manipulation Functions

- Unix provides various function calls to help you manipulate IP addresses.
- These functions convert Internet addresses between ASCII strings (what humans prefer to use) and network byte ordered binary values (values that are stored in socket address structures).
- The following three function calls are used for IPv4 addressing –
 - **int inet_aton(const char *strptr, struct in_addr *addrptr)**
 - **in_addr_t inet_addr(const char *strptr)**
 - **char *inet_ntoa(struct in_addr inaddr)**

Byte Manipulation Functions

int inet_aton(const char *strptr, struct in_addr *addrptr)

- This function call converts the specified string in the Internet standard dot notation to a network address, and stores the address in the structure provided.
- The converted address will be in Network Byte Order (bytes ordered from left to right). It returns 1 if the string was valid and 0 on error.
- example –

In the code, the `memset()` function is used to set the memory pointed to by `addrptr` to all zeros (`\0`) 1. This is often done to initialize a structure or buffer before storing data in it.

```
#include <arpa/inet.h>
(...)
int retval;
struct in_addr addrptr;
memset(&addrptr, '\0', sizeof(addrptr));
retval = inet_aton("68.178.157.132", &addrptr);
(...)
```


Byte Manipulation Functions

in_addr_t inet_addr(const char *strptr)

- This function call converts the specified string in the Internet standard dot notation to an integer value suitable for use as an Internet address.
- The converted address will be in Network Byte Order (bytes ordered from left to right). It returns a 32-bit binary network byte ordered IPv4 address and INADDR_NONE on error.
- example –

```
#include <arpa/inet.h>
(...)
struct sockaddr_in dest;
memset(&dest, '\0', sizeof(dest));
dest.sin_addr.s_addr
    =inet_addr("69.172.156.131");
(...)
```

Byte Manipulation Functions

char *inet_ntoa(struct in_addr inaddr)

- This function call converts the specified Internet host address to a string in the Internet standard dot notation.
- Example

```
#include <arpa/inet.h>
(...)

char *ip;
ip = inet_ntoa(dest.sin_addr);
printf("IP Address is: %s\n", ip);

(...)
```

IPv6 Address Structure

```
struct in6_addr {  
    uint8_t    s6_addr[16];           /* 128-bit IPv6 address */  
                                         /* network byte ordered */  
};  
  
#define SIN6_LEN        /* required for compile-time tests */  
  
struct sockaddr_in6 {  
    uint8_t      sin6_len;             /* length of this struct (28) */  
    sa_family_t  sin6_family;         /* AF_INET6 */  
    in_port_t    sin6_port;           /* transport layer port# */  
                                         /* network byte ordered */  
    uint32_t     sin6_flowinfo;       /* flow information, undefined */  
    struct in6_addr sin6_addr;         /* IPv6 address */  
                                         /* network byte ordered */  
    uint32_t     sin6_scope_id;       /* set of interfaces for a scope */  
};
```

IPv6 Address Structure

- The `SIN6_LEN` constant must be defined if the system supports the length member for socket address structures.
- The IPv6 family is `AF_INET6`, whereas the IPv4 family is `AF_INET`.
- The members in this structure are ordered so that if the `sockaddr_in6` structure is 64-bit aligned, so is the 128-bit `sin6_addr` member. On some 64-bit processors, data accesses of 64-bit values are optimized if stored on a 64-bit boundary.
- The `sin6_flowinfo` member is divided into two fields:
 - The low-order 20 bits are the flow label
 - The high-order 12 bits are reserved
- The use of the flow label field is still a research topic.

IPv6 Address Structure

IPv4




Figure 3.1

IPv6

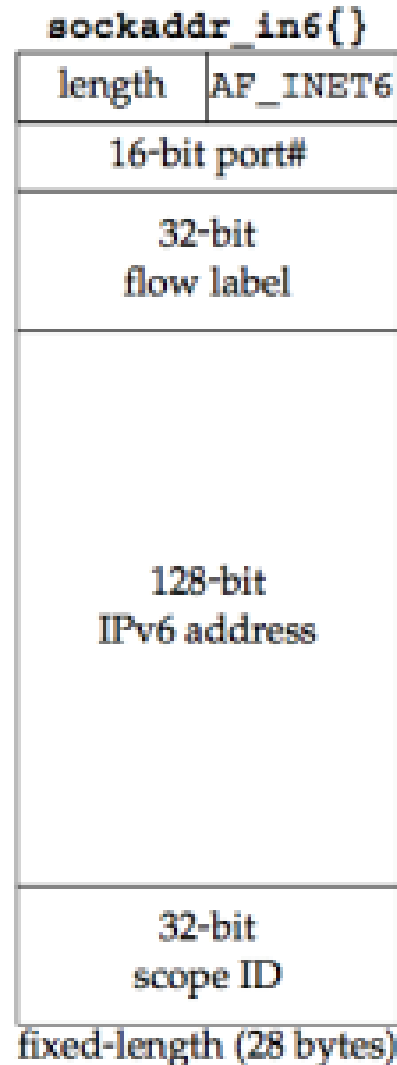


Figure 3.4

Compare IPv4 and IPv6
Address Structures

IPv6 Address Structure

Example:

```
int listen_sock_fd;  
struct sockaddr_in6 server_addr; // IPv6 address  
  
/* Create socket for listening (client requests) */  
    listen_sock_fd = socket(AF_INET6, SOCK_STREAM, IPPROTO_TCP);  
  
/* Assign Values to IPv6 address*/  
  
    server_addr.sin6_family = AF_INET6;  
    server_addr.sin6_addr = in6addr_any; // #  
    server_addr.sin6_port = htons(SERVER_PORT);  
  
    /* Bind address and socket together */  
bind(listen_sock_fd, (struct sockaddr*)&server_addr, sizeof(server_addr));  
  
/*rest all code is same as of IPv4*/  
  
# instead following may be used to assign loopback address only  
inet_pton(AF_INET6, "::1", &server_addr.sin6_addr);
```